

Sample AI Detection Tool Risk Report

AI system assessed

Field	Value
System type	AI detection / academic integrity tool
Deployment category	Current deployment risk
Student impact area	Academic misconduct
Risk tier	High
Primary concern	False positive and false negative risk
Required control	Human review and appeal route

Key audit questions

1. Does the tool flag students differently across gender groups?
2. Are false positive rates materially different by gender?
3. Are false negative rates materially different by gender?
4. Are flagged students receiving human review?
5. Are appeal outcomes monitored?
6. Is there enough evidence to justify continued use or scaling?

Required fields in real deployment logs

A university partner would need to provide anonymised logs including:

Field	Purpose
Student group attribute	Enables subgroup fairness audit
Detector score	Allows threshold and distribution testing
Detection flag	Shows who was flagged
Confirmed AI use or adjudicated outcome	Enables false positive / false negative testing
Human review status	Checks oversight quality

Misconduct case outcome	Shows operational impact
Appeal submission and appeal outcome	Tests student rights and correction mechanisms

Metrics calculated

Metric	Why it matters
Detection flag rate	Shows how often students are flagged
False positive rate	Measures wrongful suspicion risk
False negative rate	Measures missed AI-use risk
Human review rate	Tests whether oversight is meaningful
Misconduct case rate	Shows escalation impact
Appeal success rate	Indicates whether errors are corrected

Interpretation framework

A high false positive rate is serious because AI detection tools can influence academic misconduct processes and student records.

A gender gap in false positive rate does not automatically prove discrimination. It indicates that the system requires deeper review before deployment or scaling.

The correct governance response is not automatic rejection. The correct response depends on:

- Size of the fairness gap
- Reliability of the ground truth label
- Availability of human review
- Strength of student appeal rights
- Whether the tool is advisory or determinative
- Whether the vendor can provide validation evidence

Example deployment recommendation

Recommendation Rationale

Pilot only / pause deployment Use should remain limited until false positive testing, human review, appeal evidence, and subgroup monitoring are complete

Required mitigations

1. Do not allow sole reliance on AI detector output.
2. Require human review before academic misconduct escalation.
3. Provide plain-language student notice.
4. Provide a clear appeal route.
5. Monitor false positives and false negatives by gender and other available protected characteristics.
6. Retest after major vendor updates or threshold changes.
7. Keep an incident log and report unresolved disparity to the governance committee.