

Sample Governance Evidence Pack Report

Report status

This report demonstrates how the toolkit checks whether a university AI system has enough governance evidence before approval or scaling.

It is based on the project framework's evidence-pack logic and is intended for public GitHub demonstration.

Required evidence items

Evidence item	Why it matters
System description	Defines what the system is and what it does
Decision influence map	Shows where AI affects student outcomes
Data dictionary	Explains variables and data sources
Demographic coverage summary	Checks whether groups are represented
Fairness and performance test results	Shows whether the system was tested properly
Human oversight procedure	Prevents rubber-stamping and automation bias
Vendor model card	Improves transparency
DPIA summary	Supports data protection review
Equality impact assessment	Supports equality-risk review
Student notice and explanation text	Supports transparency
Appeal route	Protects student rights
Rollback or suspension plan	Allows safe withdrawal
Monitoring schedule	Supports ongoing review
Incident log	Captures harm, errors, and complaints

Evidence status scoring

Status

Provided Partial Missing **Scores:** 1.0, 0.5, 0.0

The toolkit converts evidence completeness into a control maturity score.

Governance interpretation

Finding	Recommended response
Most evidence provided	Approve or approve with monitoring
Important evidence partial	Pilot only or conditional approval
Critical evidence missing	Pause deployment
No appeal route or human oversight	Pause or reject
No fairness testing for high-risk system	Pause deployment

Critical blocker examples

A system should not be approved if:

- It affects students materially but has no appeal route.
- It flags, ranks, or profiles students without human review.
- It lacks fairness testing for protected or monitored groups.
- It has no monitoring schedule.
- Vendor documentation is unavailable.
- Students are not told when AI materially affects them.